

# TensorLayer: A Versatile Library for Deep Learning Development

**Hao Dong, Luo Mai, Simiao Yu**

Imperial College London

Acknowledge: Akara Supratak, Guo Li, Fangde Liu, Axel Oehmichen, Yike Guo

# Outline

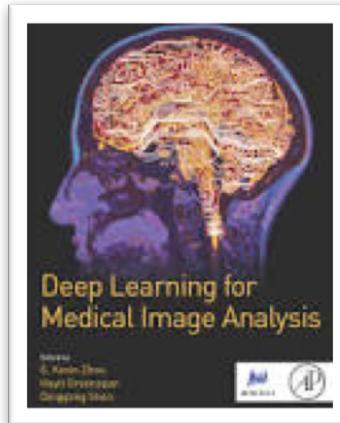
- Motivation
- Design
- Impact
- Example & Research
- Future

# Motivation

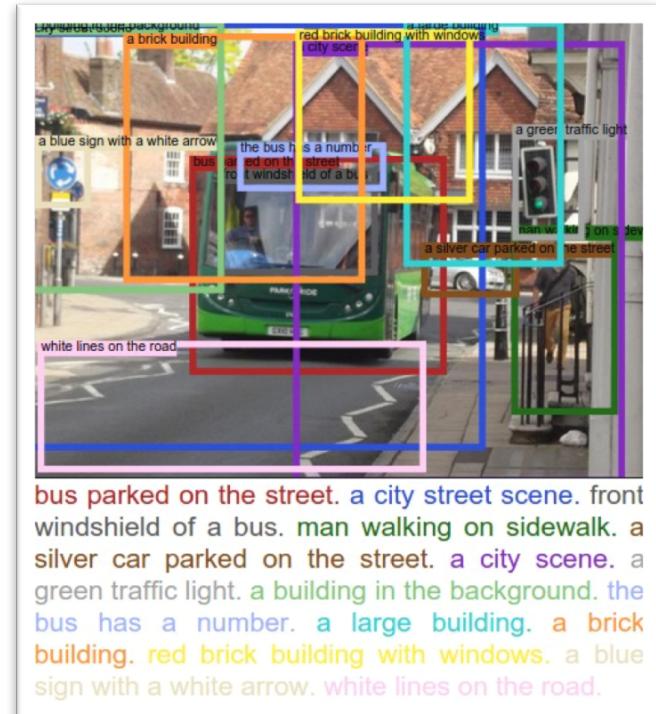
# Motivation

Deep learning has enabled many AI systems to exhibit unprecedented performance even beyond human

- Computer vision
- Natural language processing
- Medical imaging and diagnosis
- Game playing
- ...



<http://cs.stanford.edu/people/karpathy/densecap/eg/p1.jpeg>



# Motivation

## Key reasons for deep learning success

- Big data and systems
  - ImageNet, Youtube8M, Hadoop Ecosystem
- Algorithm breakthroughs
  - AlexNet, GoogleNet, AlphaGo, AlphaZero
- Affordable commodity hardware
  - GPU, TPU, FPGA and ASIC
- Fast innovations through open-source
  - Data, algorithm and hardware can be accessed by anyone

# Motivation

Open-source deep learning platforms



~90000 GitHub stars

Key reasons for TensorFlow

- Largest user base
- Widest production adoption
- Well-maintained documents
- Battlefield-proof quality



~12000 GitHub stars



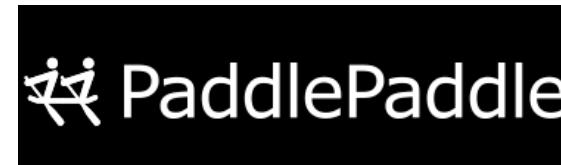
~14000 GitHub stars



~7000 GitHub stars



~12000 GitHub stars

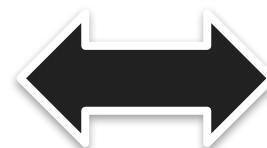
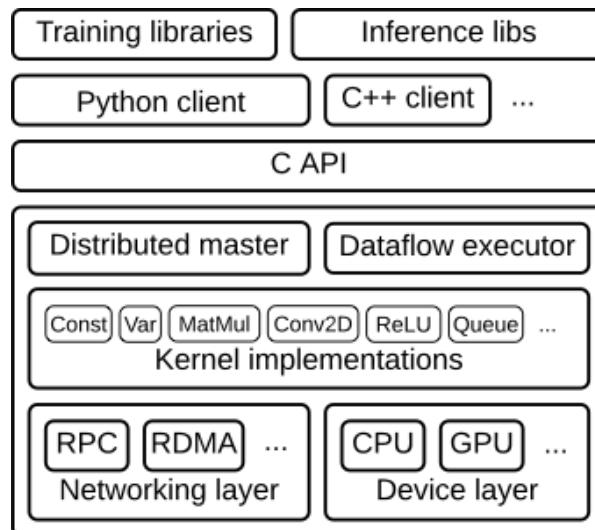


~6000 GitHub stars

# Motivation

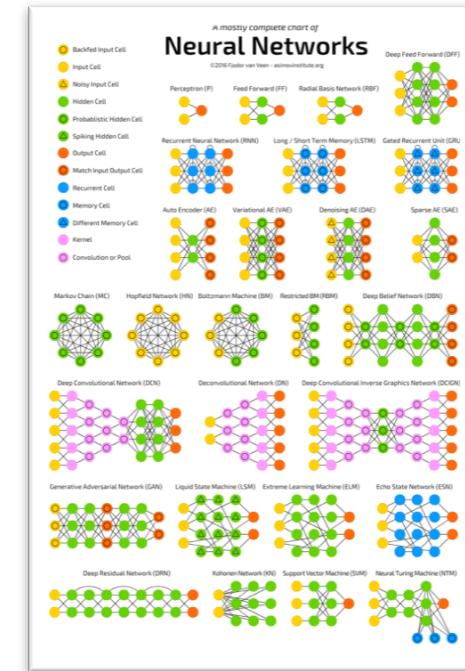
Google TensorFlow is powerful but hard to master

TensorFlow low-level interface:  
dataflow graph, placeholder,  
session, queue runner, devices ...



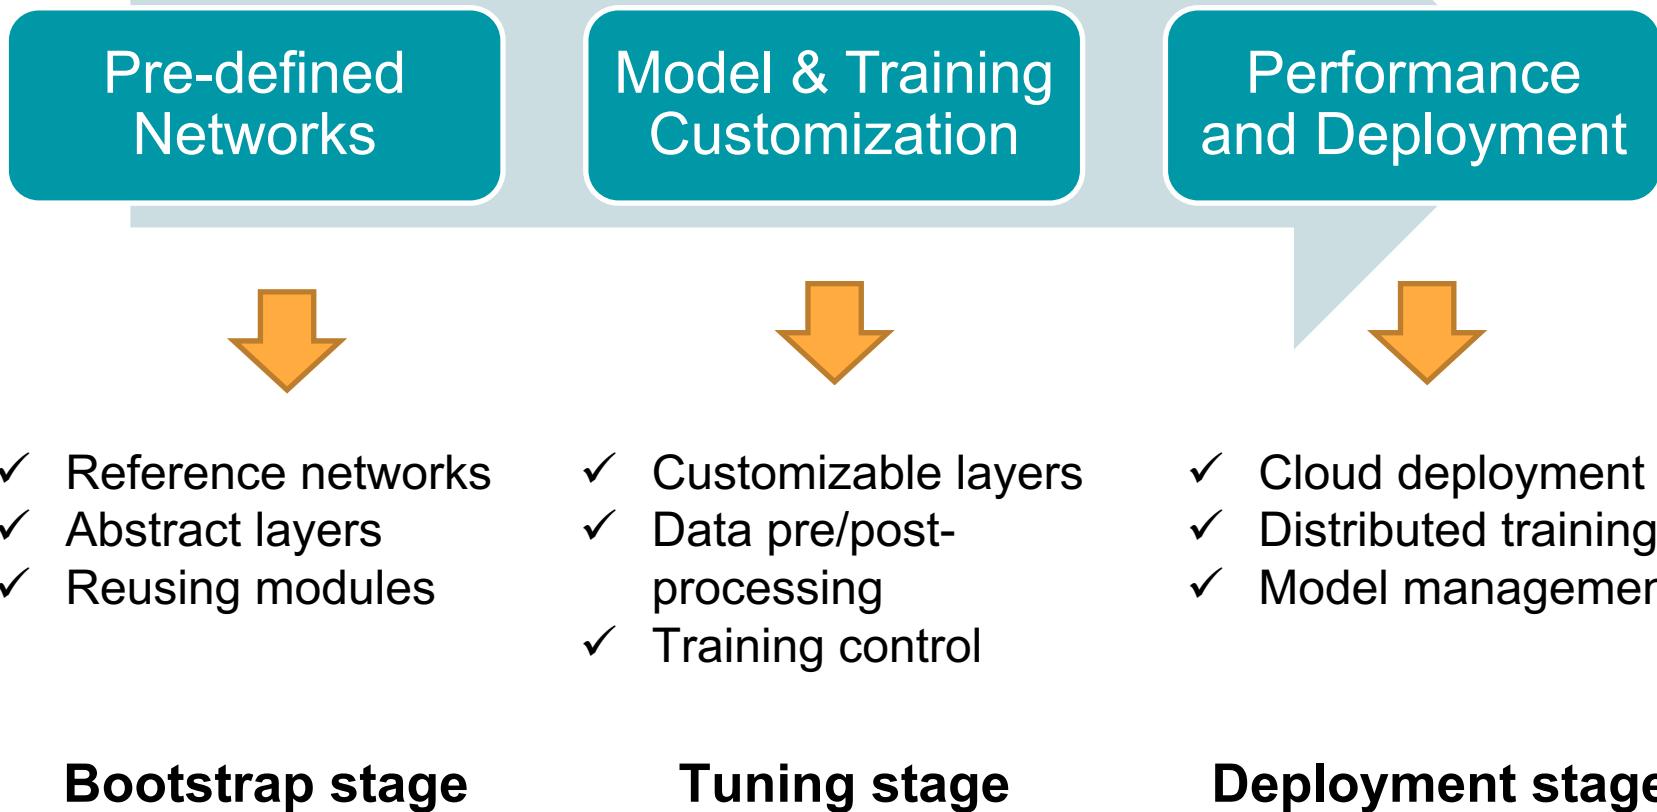
Abstraction gap

Bridged by wrappers:  
TensorLayer, Keras  
and TFLearn



# Motivation

## Library requirements during development stages



# Motivation

## Comparing TensorLayer, Keras and TFLearn

	TensorLayer	Keras	TFLearn
Abstraction	Modest	High	Modest
Reference layers	Rich	Limited	Limited
Portable backend	No	Yes	No
Flexibility	Excellent	Limited	Good
Performance & deployment	Excellent	Limited	Good

TensorLayer stands at a unique spot for its **simplicity**, **flexibility** and **performance**.

# Design

# Design Highlights

## Simplicity

- High-level state-of-the-art deep learning layers
- Mass examples and tutorials
- Mass helper functions for data preparation, augmentation, visualization and evaluation
- Import modules from Keras and TFSlim

## Transparency

- Easy to access TensorFlow operations
- Easy to insert customized model and training logics

## Performance

- Zero performance downgrade compared to raw TensorFlow
- Easy to scale from laptops to clouds

# Design Principles - Simplicity

- Provide high-level state-of-the-art deep learning modules

```

def model(x, reuse=False, is_train=False):
    with tf.variable_scope("model", reuse=reuse):
        tl.layers.set_name_reuse(reuse)
        net = InputLayer(x, name='input')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn1')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch1')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool1')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn2')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch2')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool2')

        net = FlattenLayer(net, name='flatten')
        net = DenseLayer(net, 384, act=tf.nn.relu, name='d1')
        net = DenseLayer(net, 192, act=tf.nn.relu, name='d2')
        net = DenseLayer(net, 10, name='output')
    return net
  
```

DeformableConv2dLayer

MultiplexerLayer

SubpixelConv1d

SlimNetsLayer

SpatialTransformer2dAffineLayer

SubpixelConv2d

ConvLSTMLayer

DropconnectDenseLayer

Seq2Seq

ROIPoolingLayer

UnStackLayer

# Design Principles - Transparency

- Enable users to build a model using native Tensorflow APIs

```

def model(x, reuse=False, is_train=False):
    with tf.variable_scope("model", reuse=reuse):
        tl.layers.set_name_reuse(reuse)
        net = InputLayer(x, name='input')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn1')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch1')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool1')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn2')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch2')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool2')

        net = FlattenLayer(net, name='flatten')
        net = DenseLayer(net, 384, act= tf.nn.relu, name='d1')
        net = DenseLayer(net, 192, act= tf.nn.relu, name='d2')
        net = DenseLayer(net, 10, name='output') ▼

    return net
  
```

**Pass by function**

# Design Principles - Transparency

- Enable users to build a model using native Tensorflow APIs

```

def model(x, reuse=False, is_train=False):
    with tf.variable_scope("model", reuse=reuse):
        tl.layers.set_name_reuse(reuse)
        net = InputLayer(x, name='input')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn1')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch1')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool1')

        net = Conv2d(net, 64, (5, 5), (1, 1), padding='SAME', name='cnn2')
        net = BatchNormLayer(net, is_train, act=tf.nn.relu, name='batch2')
        net = MaxPool2d(net, (3, 3), (2, 2), padding='SAME', name='pool2')

        net = FlattenLayer(net, name='flatten')
        net = DenseLayer(net, 384, act=tf.nn.relu, name='d1')
        net = DenseLayer(net, 192, act=tf.nn.relu, name='d2')
        net = DenseLayer(net, 10, name='output')

    return net
  
```

**Simplified-version API**

```

net = Conv2dLayer(net, act=tf.identity, shape=[5, 5, 3, 64],
                  strides=[1, 1, 1, 1], padding='SAME', name='cnn1')

net = PoolLayer(net, ksize=[1, 3, 3, 1], strides=[1, 2, 2, 1],
                padding='SAME', pool=tf.nn.max_pool, name='pool1')
  
```

**Advanced-version API**

# Design Principles - Transparency

- Enable users to define their own computational operations

```
net = DenseLayer(net, 200, act=lambda x: tl.act.lrelu(x, 0.2), name='d1')
```

```
net = DenseLayer(net, 200, act=lambda x: 5*x, name='d2')
```

**Lambda expression**

```
net = ElementwiseLayer([net1, net2, net3], tf.add, name='ele')
```

**tf.minimum, tf.maximum, tf.multiply, lambda x, y : x+2\*y and etc**

```
net = DynamicRNNLayer(net,
                      cell_fn=tf.contrib.rnn.BasicLSTMCell,
                      n_hidden=100,
                      name='rnn')
```

**tf.contrib.rnn.(RNNCell, GRUCell, LSTMCell, LayerNormBasicLSTMCell and etc)**

# Design Principles - Transparency

- Enable users to define their own training logic

## Data Iteration Toolbox

```

for epoch in range(n_epoch):
    start_time = time.time()
    for X_train_a, y_train_a in tl.iterate.minibatches(
        X_train, y_train, batch_size, shuffle=True):
        sess.run(train_op, feed_dict={x: X_train_a, y_: y_train_a})

    if epoch + 1 == 1 or (epoch + 1) % print_freq == 0:
        print("Epoch %d of %d took %fs" % (epoch + 1,
                                             n_epoch, time.time() - start_time))
        train_loss, train_acc, n_batch = 0, 0, 0
        for X_train_a, y_train_a in tl.iterate.minibatches(
            X_train, y_train, batch_size, shuffle=True):
            err, ac = sess.run([cost_test, acc],
                               feed_dict={x: X_train_a, y_: y_train_a})
            train_loss += err; train_acc += ac; n_batch += 1
        print("  train loss: %f" % (train_loss/ n_batch))
        print("  train acc: %f" % (train_acc/ n_batch))
        val_loss, val_acc, n_batch = 0, 0, 0
        for X_val_a, y_val_a in tl.iterate.minibatches(
            X_val, y_val, batch_size, shuffle=True):
            err, ac = sess.run([cost_test, acc],
                               feed_dict={x: X_val_a, y_: y_val_a})
            val_loss += err; val_acc += ac; n_batch += 1
        print("  val loss: %f" % (val_loss/ n_batch))
        print("  val acc: %f" % (val_acc/ n_batch))
  
```

## Training Toolbox

```

tl.utils.fit(sess, network, train_op, cost, X_train, y_train, x, y_,
             acc=acc, batch_size=500, n_epoch=500, print_freq=5,
             X_val=X_val, y_val=y_val, eval_train=False)
  
```

# Design Principles - Composability

- Users can glue different modules together (e.g., connected with [TF-Slim](#) and [Keras](#)).

```
x = tf.placeholder(tf.float32, shape=[None, 784])
y_ = tf.placeholder(tf.int64, shape=[None,])
```

```
def keras_block(x):
    x = Dropout(0.8)(x)
    x = Dense(800, activation='relu')(x)
    x = Dropout(0.5)(x)
    x = Dense(800, activation='relu')(x)
    x = Dropout(0.5)(x)
    logits = Dense(10, activation='linear')(x)
    return logits
```

```
net = InputLayer(x, name='input')
net = LambdaLayer(net, fn=keras_block, name='keras')
```

Keras

```
x = tf.placeholder(tf.float32, shape=[None, 299, 299, 3])
```

```
net = tl.layers.InputLayer(x, name='input')
```

```
with slim.arg_scope(inception_v3_arg_scope()):
    net = tl.layers.SlimNetsLayer(layer=net,
                                  slim_layer=inception_v3,
                                  slim_args={
                                      'num_classes': 1001,
                                      'is_training': False
                                  },
                                  name='InceptionV3')
```

TF-Slim

TF-Slim to TensorLayer

Keras to TensorLayer

# Design Principles - Performance

- Provide **zero-cost** abstraction (negligible overhead)

	CIFAR-10	PTB LSTM	Word2Vec
TensorLayer	2528 images/s	18063 words/s	58167 words/s
TensorFlow	2530 images/s	18075 words/s	58181 words/s

Titan X Pascal GPU 2017

# Design Principles - Performance

Scaling your models from laptops to clouds

- Scenario: Train your models using multiple GPUs in **one machine**
  - Start with TensorLayer MNIST and InceptionV3 ImageNet examples to get familiar with the distributed APIs
  - Collaborate with the **Good AI Lab (ClusterOne)**
- Scenario: Train your models using a **multi-server cluster**
  - Play with the simple Jupyter Notebook that back up by public clouds including Google Cloud Engine, Amazon AWS and Microsoft Azure
  - Collaborate with the **Google Kubeflow team**

# Tutorials and Examples

Distributed Training	DDPG	Actor-Critic	RNN
SRGAN	A3C	Policy Gradient	VAE
Inception	Text generation	Word2vec	GAN-CLS
Data augmentation	Chatbot	Anti-Spam	Autoencoder
CNN	DCGAN	DAGGER	LSTM
VAE-GAN	Translation	Semantic Image Synthesis	Image-to-image translation
Deep Q Network	VGG	Brain tumor segmentation	

# Installation and Documentation

## Installation:

```
pip install tensorlayer
```

## Platform:

Linux, Windows, OSX

## Documentation:



<http://tensorlayer.readthedocs.io> or <http://tensorlayer.org>



<http://tensorlayercn.readthedocs.io>

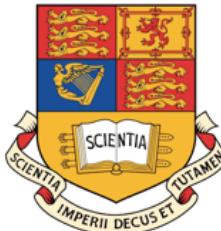


Coming soon

# Impact

# Impact - Github

Jan 2018 ~ 3000 Stars ~ 700 Forks ~ Global Contributors



清华大学  
Tsinghua University



Stanford  
University



PENGUINS  
INNOVATE



LINKÖPING  
UNIVERSITY



Tencent 腾讯



ReFUEL4

UCLA



Good AI Lab

Bloomberg

Google

Alibaba Group



小米  
xiaomi.com



Microsoft

# Impact - Media and Award



Good AI Lab

Humans of AI



TensorLayer: A Versatile Library for Efficient Deep Learning Development. *H. Dong, A. Supratak et al.* ACM MM 2017.



# Impact - Book from Community

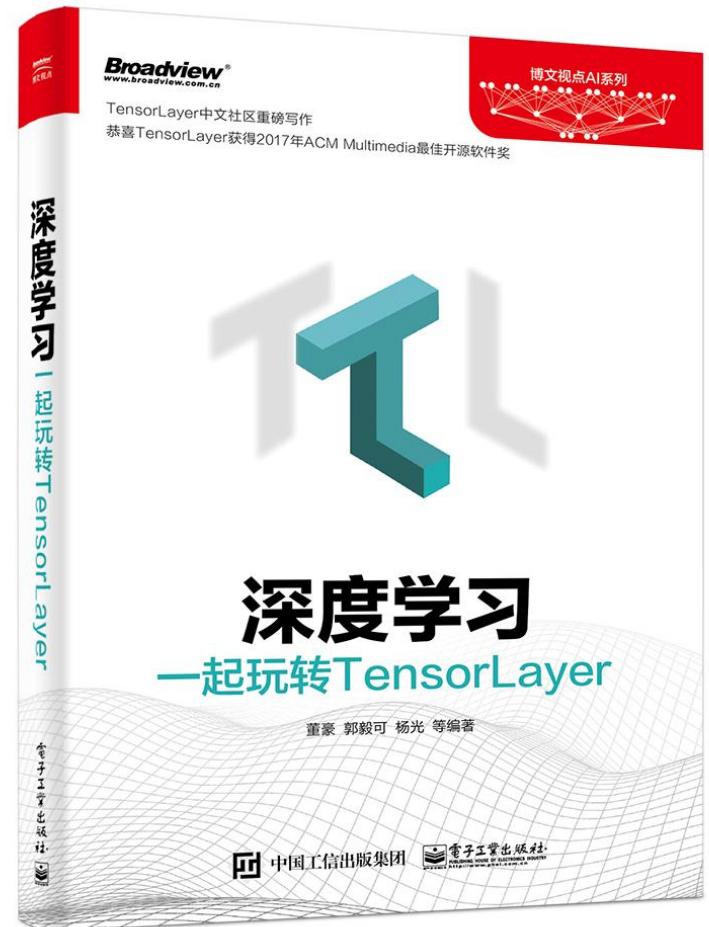
Invited by Publishing House of Electronics Industry (PHEI, 电子工业出版社)

“Deep Learning using TensorLayer”

10 contributors

Published date: Mar 2018

English version will come later



# Example & Research

# Research - Computer Vision

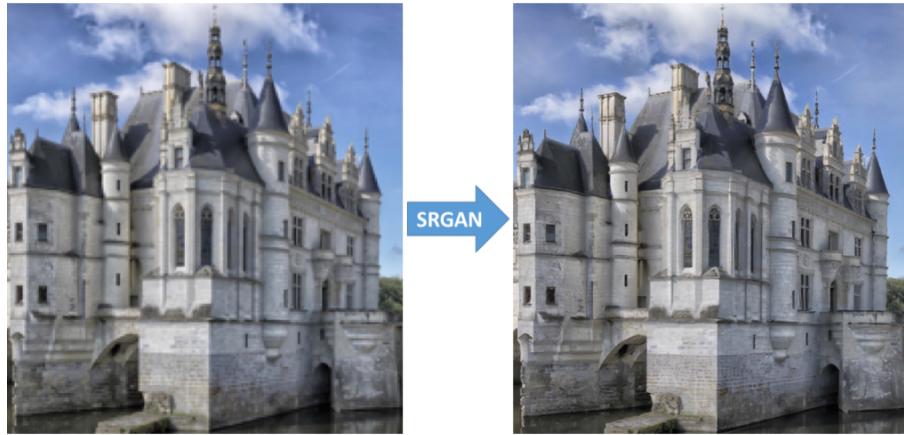
## Unsupervised Image to Image Translation



Person A to B

Person B to A

## Super Resolution



LG Image

Generated Image

Semantic Image Synthesis via Adversarial Learning. *H. Dong, S. Yu et al. ICCV 2017.*

## Semantic Image Synthesis

The bird has blue  
+ crown and wings, and = white breast.



+ A red bird with blue  
head has grey wings. =



This flower has white  
+ petals with yellow  
round stamens. =



This flower is pink and  
+ white in color, and has =  
no visible stamens.

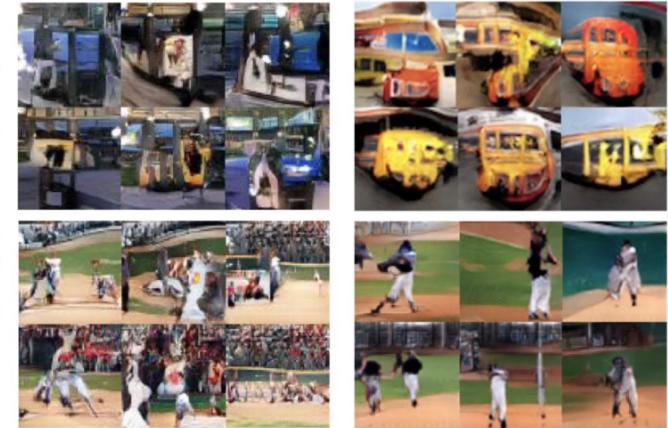


## Text to Image Synthesis

GAN-CLS

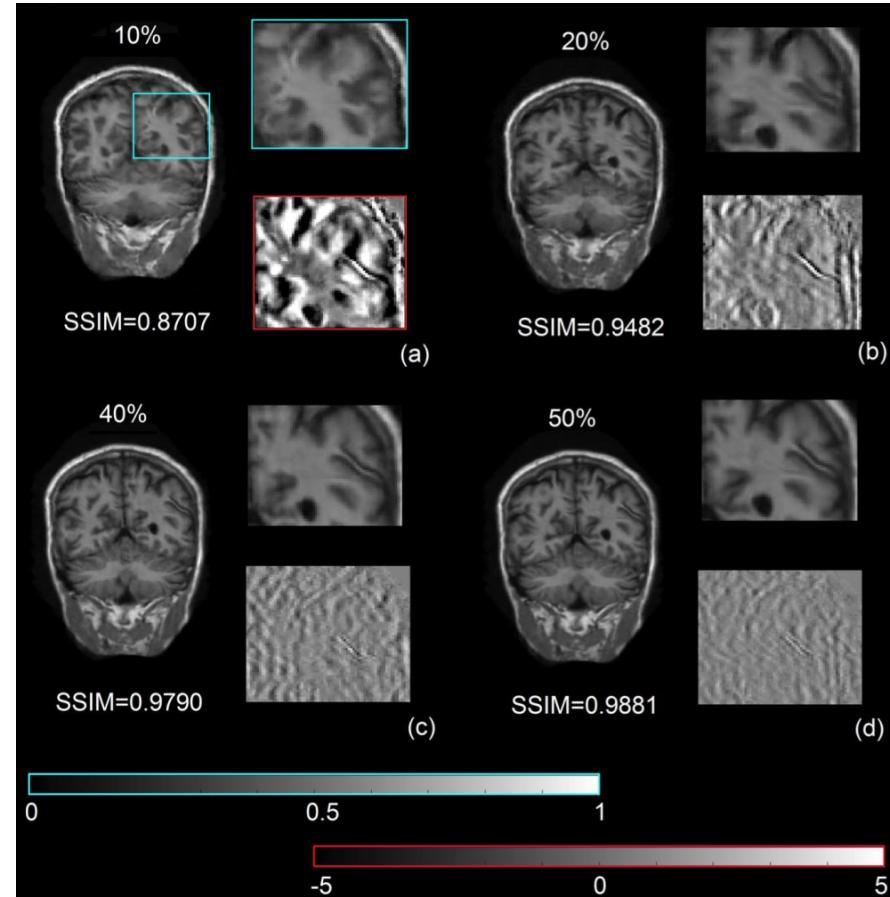
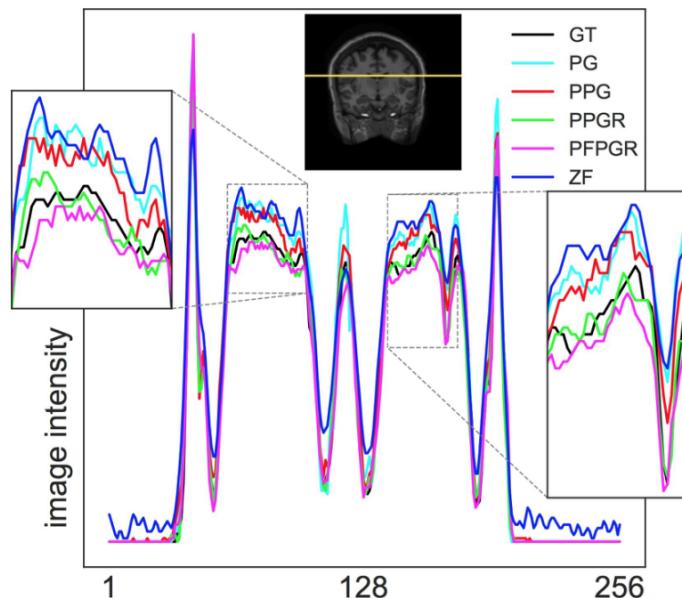
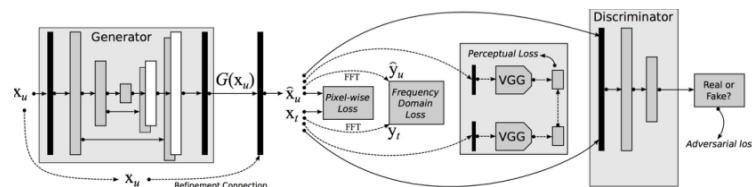


I2T2I



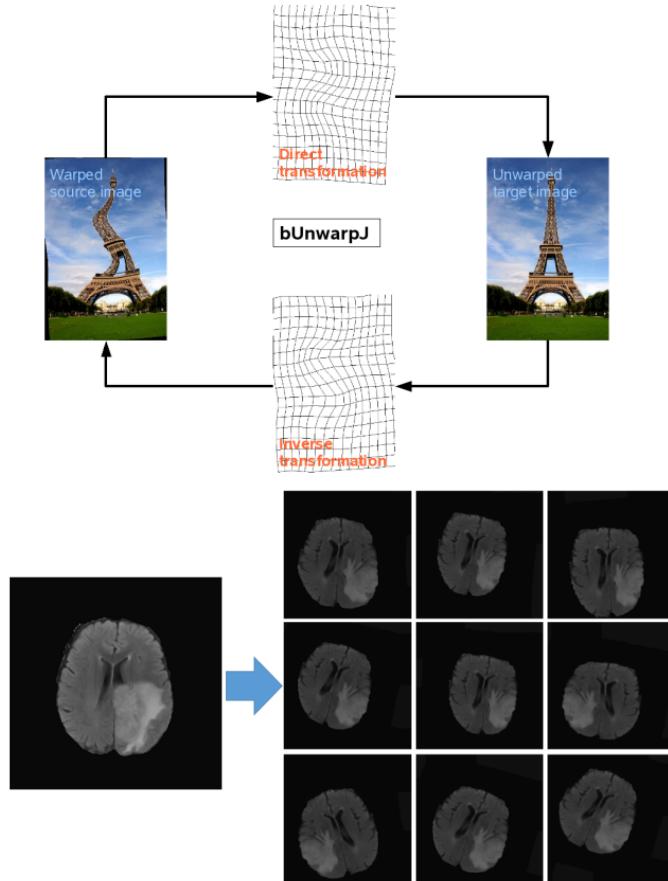
# Research - Medical Applications

## Super Resolution for Medical Image



# Research - Medical Applications

## Elastic Transformation for MRI Segmentation



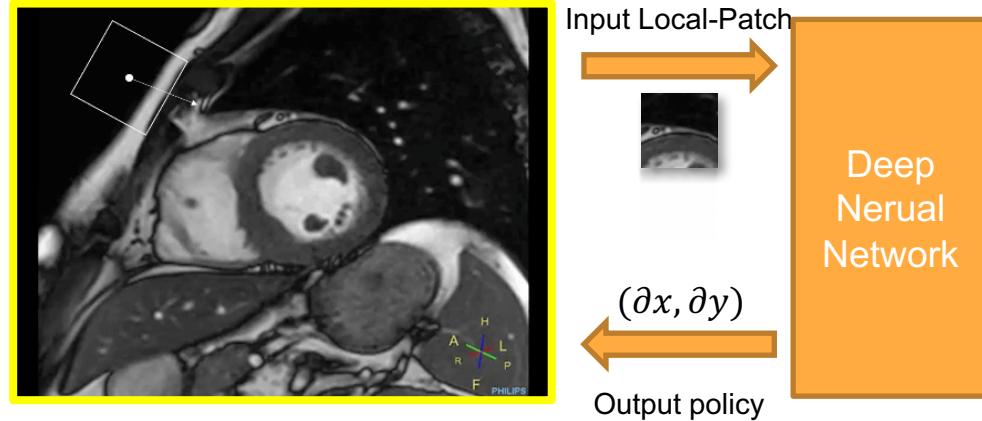
```

def distort_imgs(data):
    """ data augmentation """
    x1, x2, x3, x4, y = data
    x1, x2, x3, x4, y = tl.prepro.flip_axis_multi([x1, x2, x3, x4, y],
                                                    axis=1, is_random=True) # left right
    x1, x2, x3, x4, y = tl.prepro.elastic_transform_multi([x1, x2, x3, x4, y],
                                                          alpha=720, sigma=24, is_random=True)
    x1, x2, x3, x4, y = tl.prepro.rotation_multi([x1, x2, x3, x4, y], rg=20,
                                                   is_random=True, fill_mode='constant') # nearest, constant
    x1, x2, x3, x4, y = tl.prepro.shift_multi([x1, x2, x3, x4, y], wrg=0.10,
                                                hrg=0.10, is_random=True, fill_mode='constant')
    x1, x2, x3, x4, y = tl.prepro.shear_multi([x1, x2, x3, x4, y], 0.05,
                                                is_random=True, fill_mode='constant')
    x1, x2, x3, x4, y = tl.prepro.zoom_multi([x1, x2, x3, x4, y],
                                              zoom_range=[0.9, 1.1], is_random=True,
                                              fill_mode='constant')
    return x1, x2, x3, x4, y
  
```

# Research - Medical Applications

## Deep Poincare Map (DPM) for MRI Segmentation

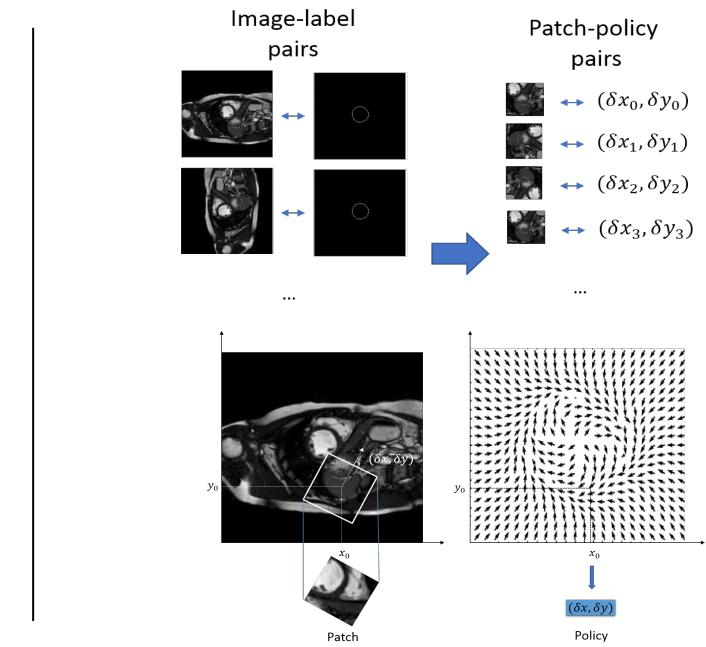
### Animation of Inference



Unlike tradition End-to-End DL method.

DPM is a iterative process to segment ROI.

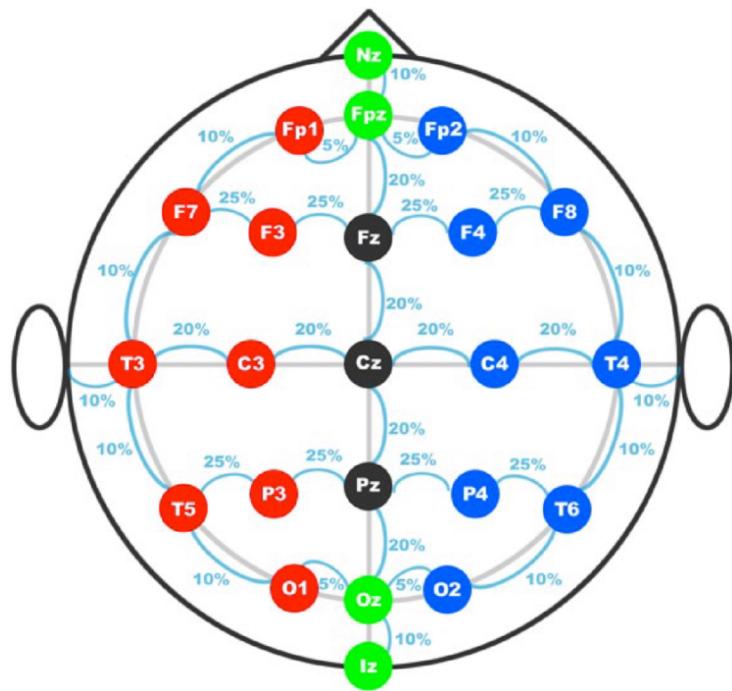
### Training



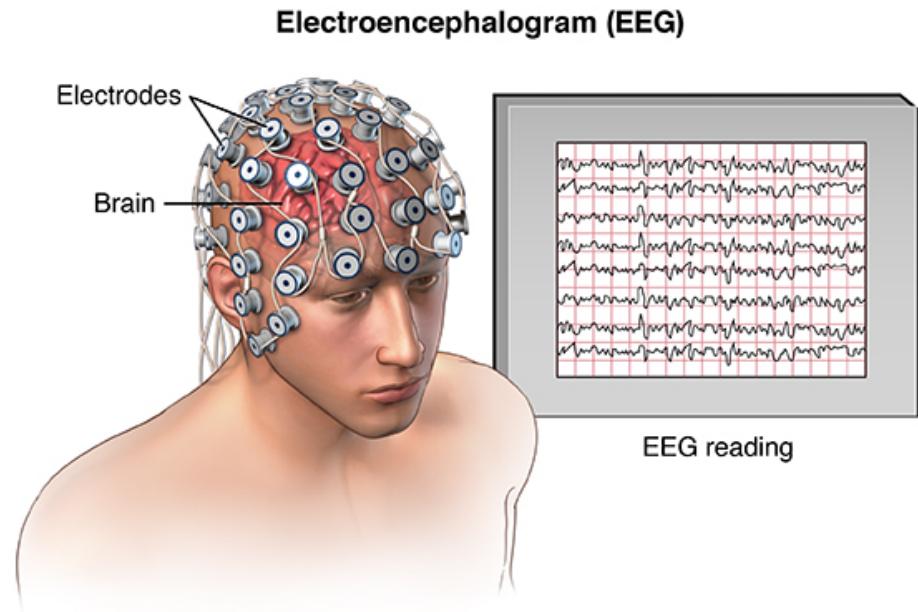
We will generate a dynamical system for each so that we can sample local patch and its corresponding policy

# Research - Medical Applications

## Wearable EEG Sleep Scoring



10/20 EEG electrode standard



Multi-channel EEG device

Mixed Neural Network Approach for Temporal Sleep Stage Classification. *H. Dong, A. Supratak et al. TNSRE 2016.*

DeepSleepNet: a Model for Automatic Sleep Stage Scoring based on Raw Single-Channel EEG.  
*A. Supratak, H. Dong et al. TNSRE 2017.*

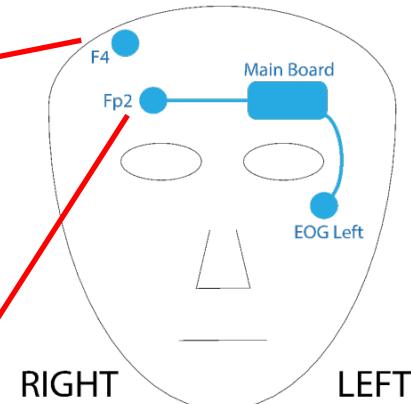
# Research - Medical Applications

## Wearable EEG Sleep Scoring

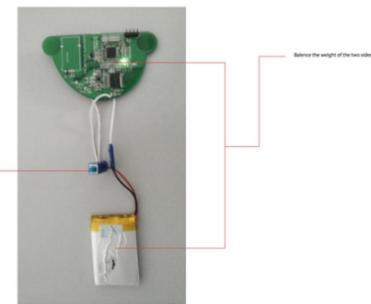
**83~86% accuracy**

CONFUSION MATRIX FROM CROSS-VALIDATION USING MNN AND F4-EOG (LEFT) WHEN SEQUENCE LENGTH IS 5. THE LEFT COLUMN IS THE ACTUAL CLASS LABELED BY SLEEP EXPERT, AND THE TOP ROW IS THE PREDICTED CLASS CALCULATED BY THE CLASSIFIER.

SL=5	W	ACC = 85.92% MF1 = 80.50%					
		N1	N2	N3	REM	RE	PR
W	5022	577	188	19	395	80.95	88.49
N1	407	2468	989	4	965	51.07	62.75
N2	130	630	27254	1021	763	91.46	90.02
N3	13	0	1236	6399	5	83.61	85.94
REM	103	258	609	0	9611	90.83	81.87
						86.12	

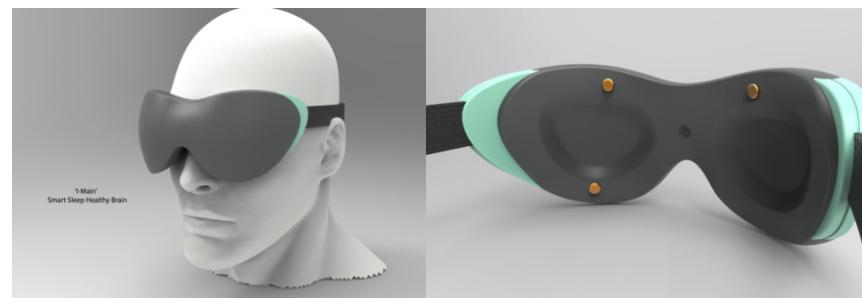


## Active Dry Electrode



CONFUSION MATRIX FROM CROSS-VALIDATION USING MNN AND *Fp2-EOG (Left)* WHEN SEQUENCE LENGTH IS 5

SL=5	W	ACC = 83.35% MF1 = 76.97%					
		N1	N2	N3	REM	RE	PR
W	4604	795	294	32	479	74.21	86.04
N1	405	2208	1292	9	919	45.69	57.08
N2	208	605	27199	897	889	91.28	86.65
N3	24	1	1689	5936	3	77.56	86.22
REM	110	259	914	11	9287	87.77	80.22
						83.83	



Mixed Neural Network Approach for Temporal Sleep Stage Classification. *H. Dong, A. Supratak et al. TNSRE 2016.*

DeepSleepNet: a Model for Automatic Sleep Stage Scoring based on Raw Single-Channel EEG.  
*A. Supratak, H. Dong et al. TNSRE 2017.*

A New Soft Material based In-the-ear EEG Recording Technique *H. Dong, PM. Matthews, Y. Guo. EMBC 2016.*

# Future

# Future

Moving from the current 1.8 to the 2.0 milestone

- Support more domain-specific layers
- Support for integrated command line interface
- Support for distributed training and large dataset
- Support for model compression



# Happy Chinese New Year

新春快乐